

See demand forming. Decide with evidence.

An earlier, auditable destination review from incomplete search interest.

321 cities

7 provinces

30 check-in dates

Two clocks. One observed-demand floor.

Final searches = observed + forecast remaining

Search date tells us what was knowable.

Check-in date tells us what we are forecasting.

Two LightGBM horizon bands · guarded remainder calibration · validated 9,630-row submission

Measured improvement. The difficult fold stays in.

14.79%


Five-fold WAPE

22.00%

Pickup baseline

10.61%

Without shock fold · diagnostic only



May 21 fold overlaps June 13–24, 2025 war; 46.30% of error. Temporal association, not causal proof. WAPE 95% fold CI: 9.76%–22.18%.

Clustering had to earn its place.

0.120168

City model
Original three-fold sweep

0.117952

Same city predictions
Summed onto merged groups

0.117699

Cluster model
On those same groups

About 90% of the apparent gain was mechanical error cancellation.

We retain city-level submission rows. The 0.06% blend gain fails the greater-than-1% practical margin.

Three-fold uncalibrated experiment: compare within this slide only, not with the five-fold headline.

One destination. A review with an owner.

Example from the historical competition snapshot

tehran

1,175,144 forecast searches · 29.16% observed / forecast

Owner: Destination growth analyst

Action: Review destination content and campaign calendar

Decision gate: Join conversion, available inventory, margin and campaign cost; run a controlled pilot

[Open the live evidence room →](#) · filter a province · select a city · export the queue

Prove the workflow. Then price the value.

Week 1

Define cohort, baseline analyst time, decision log and sponsor.

Week 2

Shadow current planning.
Measure review time and corrections.

Weeks 3-4

Evaluate a preregistered pilot.
Join bookings, inventory and costs before commercial claims.

Proposed gate: $\geq 20\%$ lower median review time, with no rise in correction rate.

Buyer hypothesis: destination growth lead. Offer hypothesis: scoped evaluation, then a monthly workspace. Pricing and willingness-to-pay are unvalidated.

Inspectable evidence. A concrete pilot ask.

- Champion trainset recovered: every original file hash matches.
- Submitted forecast preserved; challengers evaluated separately.
- Search-interest scope, shock sensitivity and model limitations disclosed.

We seek an operator-sponsored pilot with timestamped booking, inventory and campaign-cost data.

7:00 — questions welcome.

Answer first. Then show the evidence.

▼ What has been proven commercially?

Forecast accuracy and a usable review workflow. No customer traction, booking lift or realized ROI is claimed. The four-week pilot measures planning time and review quality.

▼ Can you prove no leakage?

Horizon-indexed pickup features respect the forecast cutoff. The original city-history summaries include later within-fold training outcomes. The direct 66-feature origin-history challenger scores 0.151430 calibrated WAPE versus 0.147927 for the original setup on the same folds. Adversarial tests verify earlier features ignore later outcomes; pickup curves remain fold-fitted.

▼ Why keep the shock fold?

Dropping a fold because it was difficult is outcome-driven selection. Both all-fold and without-fold scores are disclosed; only the all-fold figure is the headline.

▼ Why not use the lowest challenger score?

The specifications and cluster height are exploratory, not nested. Promotion requires a prespecified rule and fresh closed holdout; a small retrospective gain is insufficient.

▼ What does the confidence interval mean?

The wide interval quantifies uncertainty of pooled WAPE over five folds. Experimental residual bands have measured forward-only coverage and no guarantee under regime shifts.

▼ Do you model lunar holidays?

No verified year-specific Iranian lunar-event table has been supplied or verified. We do not invent dates or encode unexpected war as a known-future holiday.

▼ What is the commercial ask?

A growth lead sponsor, an analyst, timestamped outcome/inventory/cost joins, and permission to run a preregistered evaluation. Willingness-to-pay and deployment costs must be measured.

[Open full experiment table and checksums →](#)